

# Statistical Experiments 24–27 — Actual Results (Plain English)

Open the live tables canvas beside chat, or use the CSVs under ``6_Results/``.

---

## Experiment 24 — Are we over-reusing the same customers?

**What we measured:** For each class we take many random snapshots. Each snapshot has `t`` people. We took `l = 500`` snapshots. The **reuse score**  $(t \times l) / n$  is roughly how many times a typical person appears across snapshots. Target: about **1 or less**. Ours are much higher.

| Dataset | Setting | People per class | Points per snapshot (t) | Snapshots (l) | Reuse score | OK? |

|-----|-----|-----|-----|-----|-----|-----|

| Credit Card | L5 (5%) | 6,630 | 331 | 500 | **25.0** | No |

| Credit Card | L15 (15%) | 6,630 | 994 | 500 | **75.0** | No |

| German Credit | L30 (30%) | 300 | 90 | 500 | **150.0** | No |

| German Credit | L60 (60%) | 180 | 180 | 500 | **300.0** | No |

**Suggested snapshots if we keep `t`` the same:**

| Dataset | Setting | Current l | Suggested l |

|-----|-----|-----|-----|

| Credit Card | L5 | 500 | **20** |

| Credit Card | L15 | 500 | **7** |

| German Credit | L30 | 500 | **3** |

| German Credit | L60 | 500 | **2** |

Files: ``6_Results/24_Sampling_Ratio_Audit/sampling_ratio_audit.csv``, ``suggested_l_values.csv``

---

## Experiment 25 — Average barcode feature values across snapshots

**What we measured:** Each snapshot → 24 numbers. Across 1,000 snapshots we stored mean and variance of every feature.

**Highlights:**

| Dataset / setting | Feature | Mean | Variance |

|-----|-----|-----|-----|

| Credit Card L5 | Mean Death (H) | 0.193 | 0.00013 |

| Credit Card L5 | Mean Persistence (H) | 0.023 | 0.000003 |

| Credit Card L15 | Mean Death (H) | 0.144 | 0.00007 |

| Credit Card L15 | Mean Persistence (H) | 0.018 | 0.000001 |

| German Credit L30 | Mean Death (H) | 1.121 | 0.00135 |

| German Credit L30 | Mean Persistence (H) | 0.114 | 0.00015 |

| German Credit L60 | Mean Death (H) | 0.996 | 0.00035 |

| German Credit L60 | Mean Persistence (H) | 0.110 | 0.00005 |

Full 96-row table: `6\_Results/25\_Snapshot\_Mean\_Variance/snapshot\_mean\_variance.csv`

---

## Experiment 26 — How “high-dimensional” is the data?

**\*\*What we measured:\*\*** Intrinsic dimension with Two-NN (main number). Roughly: how many degrees of freedom the cloud needs. Concern from the team: if this is ~7 we may be in trouble. It is **\*\*not\*\***.

| Dataset | Two-NN (all features) | Two-NN (after PCA) | PCA components used in TDA | Variance kept |

|-----|-----:|-----:|-----:|-----:|

| Credit Card | 3.95 | **\*\*2.81\*\*** | 7 | 94.0% |

| German Credit | 5.34 | **\*\*4.06\*\*** | 15 | 89.3% |

File: `6\_Results/26\_Intrinsic\_Dimension\_Estimation/intrinsic\_dimension\_estimates.csv`

---

## Experiment 27 — Do default and non-default snapshots differ?

**\*\*What we measured:\*\*** Permutation test comparing default vs non-default snapshot clouds.

**\*\*p = 0.005\*\*** means: under 200 random label shuffles, the real split looked more separated than almost all shuffles → classes look different.

| Dataset | Test | Observed score | Mean if labels shuffled | p-value | Verdict |

|-----|-----|-----:|-----:|-----:|-----|

| Credit Card L5 | F■,■ | 0.0109 | 0.0131 | 0.005 | Differ |

| Credit Card L5 | F■,■ | 0.2200 | 0.2888 | 0.005 | Differ |

| Credit Card L15 | F■,■ | 0.0069 | 0.0089 | 0.005 | Differ |

| Credit Card L15 | F■,■ | 0.1611 | 0.2364 | 0.005 | Differ |

| German Credit L30 | F■,■ | 0.1015 | 0.1139 | 0.005 | Differ |

| German Credit L30 | F■,■ | 1.0151 | 1.0662 | 0.005 | Differ |

| German Credit L60 | F■,■ | 0.0454 | 0.0560 | 0.005 | Differ |

| German Credit L60 | F■,■ | 0.6310 | 0.6995 | 0.005 | Differ |

File: `6\_Results/27\_Null\_Hypothesis\_Algorithm2/algorithm2\_permutation\_results.csv`

---

## Bottom line

1. **\*\*500 snapshots is too many\*\*** — reuse scores are 25×–300×; suggested `l` is 2–20.
2. **\*\*Barcode features are stable\*\*** — means logged; variances small.
3. **\*\*Intrinsic dimension (Two-NN) is below 7\*\*** on both datasets.
4. **\*\*Default vs non-default snapshots differ\*\*** on every test we ran (p = 0.005).